

# Meeting Minutes - Meeting-2025-11-06

Date: 2025-11-06

## Summary

Hi, this is Eric Johnson. It's February 18, 2021, and this is the Engineering Key Review at GitLab. I've got number four in the agenda, which is a proposal to break up this meeting into four department key reviews.

## Key Points

And then ignoring the seasonal dip here, we raised it to, I think, 11 and then we realized we're kind of hitting that point of diminishing returns and the thing to do business-wise, and this is in our FY 22 direction, is hold the line at productivity, but start to raise other things related to quality, security, availability, and whatnot.

I have the reasons why, increase visibility, able to go deeper, increase the objectivity with which my reports can manage their groups, allow me more time to focus on new markets, and allow me to shift into more of a question asker mode than generating content and answering questions in these meetings.

One of the most specific actions we're going to take though is separating out and having a dedicated host that we're just dealing with the profile of the data engineering traffic on there and not having conflicting query, conflicting queries affect the ability to update the replication.

, if you compare October to October, November to November and January, I'm sorry, November, December and January, comparatively to last year, which you'll find is we're between a half point and a 1.5 MR rate above where we were in the month of previous year.

then I think part of my critique of this can be thrown out because it's not as duplicative as I thought, but I still think there's a problem with R&D wider MR rate, which is this thing doesn't really move in part because it's a rate.

## Action Items

The wider MR rate should just be external, ?

And an overall should be narrow plus wider.

Infrastructure will oftentimes contribute to developments work, but it won't be counted as MRs.

That's a potential bit of funkiness that we should talk about separately.

That should not be the thing.

It should be contributions per GitLab team member.

Unless we start calling people outside the company team member, then it shouldn't be done.

The context is as we grow as a company, we should make sure we keep the community up.

If you as our CTO don't even understand them, we went overboard.

And I know we have an action plan for that as far as creating another dedicated host just for the data team to pull from.

One of the most specific actions we're going to take though is separating out and having a dedicated host that we're just dealing with the profile of the data engineering traffic on there and not having conflicting query, conflicting queries affect the ability to update the replication.

And I think we need to get, I'm still unclear on where to focus and to get the biggest bang for the buck.

I think if the computational resource dedication, that's going to be a good thing, but I'm probably going to squeeze the balloon and then the next area will unearth itself.

And this will likely be the case.

If the age goes up or down, it would be cleaning the backlog, the average age should go down as well.

We need to dig in a bit deeper in that.

, I think we shouldn't be looking at the closed bugs.

I think we should be looking at open bugs.

, you have a chart that should go up and to the , everything else.

You'll have to figure out how to, because we to be able to have charts that we can historically reconstruct if we need to.

, when tickets close out, you need to go through their history to figure out at this time when it was open, did it breach the SLO or not?

Is there anything the team wants to call out supposed to that we should be watching?

If we don't and we flatline and address it, that we flatline open S1s and S2s, you will see a temporary jump in above SLOs as we clear out that backlog over that period of time.

We need to get more security work prioritized or hearing that from the team, but neither that problem nor that activity is currently reflected in our security metrics.

, we have some work to do long-term to make sure we see things that in the metrics and the measurements that we're making.

, it should stair step here and we could go back and reconstruct that if we really wanted to.

And then ignoring the seasonal dip here, we raised it to, I think, 11 and then we realized we're kind of hitting that point of diminishing returns and the thing to do business-wise

, and this is in our FY 22 direction, is hold the line at productivity, but start to raise other things related to quality, security, availability, and whatnot.

, that's kind of why you're seeing this bump as we raised it and we realized, , that's not the, we shouldn't raise it anymore and we brought it back down to 10.

, 10 will be the target going forward.

And I just want to call out that it's not necessary, higher MR rates should also help to address security and quality and other things because you're more productive, you can fix more things.

## Deadlines

They would be viewed community contributions by that group, but effectively they're not from outside the company.

, we only went down by a tenth of a point.

## Transcript (begin)

Hi, this is Eric Johnson. It's February 18, 2021, and this is the Engineering Key Review at GitLab. I've got number four in the agenda, which is a proposal to break up this meeting into four department key reviews. currently this is Engineering, Development, Quality, Security, and UX infrastructure and support do their own key reviews already. I have the reasons why, increase visibility, able to go deeper, increase the objectivity with which my reports can manage their groups, allow me more time to focus on new markets, and allow me to shift into more of a question asker mode than generating content and answering questions in these meetings. But to avoid adding three net new meetings to stakeholders' calendars, I propose we do a sort of two-month rotation. month one, Development, Quality go, month two, Security and UX would go. How do people feel about that proposal? I think in the group conversations, it's working really well. I'm supportive. And this is the smallest thing maybe we need for meetings a month, it's the biggest department, it's super essential. But you proposed this, I don't, I could see either way. let's stick with the proposal. Cool, we'll try it and we can be flexible. , the development is larger, maybe they go more frequently or something that, but we'll see how it goes. All , and then I've got number five, which is we've got R&D overall MR rate, and we also have R&D wider MR rate, both as top level KPIs for engineering. the difference between them in the simplest sense is that R&D wider MR rate includes both community contributions and community MRs. The problems I see with

this are that one, the wider MR rate, the one that includes internal and external MRs, it duplicates the overall MR rate, which is a productivity metric. Sorry, sorry, sorry. The wider MR rate should just be external, ? And an overall should be narrow plus wider. It's we say the wider community. , , . . there's, I have to check the taxonomy. Lily, can you confirm Sid's reasoning as well? Yeah, I believe wider MR rate just captures community contributions. Only and no internal. Yeah, and the reason we measured that is that one of the most likely failure modes is that we lose the community. Yeah. where it gets goofy is that when you look at a specific team within the company, there could be contributions outside of that that aren't community contributions. They would be viewed community contributions by that group, but effectively they're not from outside the company. that's why we use wider to kind of reflect that. And narrow is very specific to the team. Are you saying that if someone in plan contributes to Verify, it's viewed as wider? Not quite that. Plan and Verify are just fine. It's when you look at the development versus infrastructure. Infrastructure will oftentimes contribute to developments work, but it won't be counted as MRs. . That's a potential bit of funkiness that we should talk about separately. I didn't have that in my sort of critique of this, but that doesn't necessarily make intuitive sense to me . then I think part of my critique of this can be thrown out because it's not as duplicative as I thought, but I still think there's a problem with R&D wider MR rate, which is this thing doesn't really move in part because it's a rate. it feels the way to drive this up is to specifically drive community authors to contribute more than one MR per month. That's how this moves up because it's a productivity rate we use internally. And that doesn't necessarily feel the thing because there's scenarios in which this goes up. We've actually got less contributions overall and less contributors overall. Wait, wait, wait, wait a second. you're saying that R&D wider MR rate is number MRs per external contributor. Oh my goodness. That should not be the thing. It should be contributions per GitLab team member. the thing above the division is the external ones. The thing below is the number of team members at GitLab. Is that the case, Lily?...